

Services and Processes

Lecture 4

aly, lemurseven

(slide credits kmo, rrchan, cooper, keur)

Course Resources

- ▣ Your facilitators!
- ▣ Ed, Gradescope
- ▣ OCF Slack (ocf.io/slack) or Discord (ocf.io/discord)
#decal-general
- ▣ All materials available at decal.ocf.io
- ▣ Ask questions / work on lab with us during lab sessions!
(Tuesday 8-9pm in OCF Lab)



Outline

- Processes
- Services
- Intro to systemd



Disclaimer

- ▣ This is not cs162
- ▣ This topic gets very deeply technical and we will only scratch the surface here
- ▣ *Main goal:* become a more productive Linux user, not a kernel hacker



Processes

What is a process?

A process is a **single instance of a program**.

Processes are isolated from one another and have their own memory, threads, etc. (Additional isolation, such as filesystem or network isolation, is also possible.)



What is a process?

- PID: Process ID
 - PPID: Parent's PID
 - UID: User running the process
 - The program (executable) that the process is running
 - The args (command line) of the process
- (and more...)



Init

- ▣ First process started at boot, given PID 1
 - ▣ Manages all other services and processes
- ▣ Run `htop`, open tree view (f5). What is the root of the tree?

PID	Command
1	/lib/systemd/systemd --system --deserialize 19
24854	└─ /usr/sbin/sshd -D
6392	└─ sshd: rrchan [priv]
6411	└─ sshd: rrchan@pts/0
6413	└─ -bash
7458	└─ htop
23491	─ nginx: master process /usr/sbin/nginx -g daemon on; master_process on;
23496	└─ nginx: worker process
23494	└─ nginx: worker process
21318	─ /lib/systemd/systemd-udevd
9077	─ tmux
11962	└─ -bash
9118	└─ -bash
9097	└─ -bash
9078	└─ -bash

Process Hierarchy

Each process is created
by a parent

(Except PID 1)

Processes can have
many children

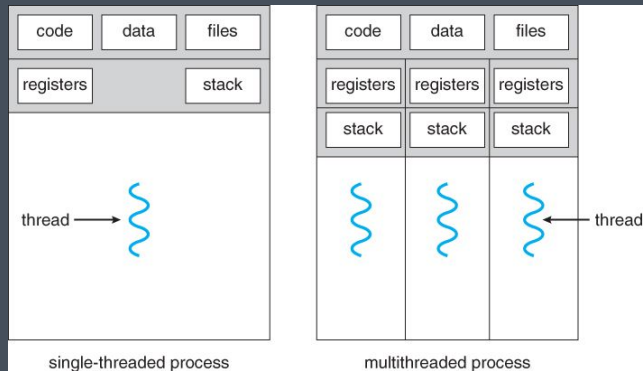
```
Command
/sbin/init
├── /lib/systemd/systemd --user
│   └── (sd-pam)
├── tmux
│   └── -zsh
│       └── python run.py
│           └── python slackbot.py
│               ├── python slackbot.py
│               ├── python slackbot.py
│               ├── python slackbot.py
│               └── python slackbot.py
```



Processes vs Threads

A process has one or more threads.

- Processes have their own data and code, must use pipes, files, etc. to communicate with one another
- Threads share the same process but have different system states (“multithreaded process”)



Aside: Why does Chrome spawn so many processes?

Command

```
/opt/google/chrome/chrome
├── /opt/google/chrome/chrome --type=gpu-process --field-trial-handle=59325916
│   └── /opt/google/chrome/chrome --type=-broker
├── /opt/google/chrome/chrome-sandbox /opt/google/chrome/chrome --type=zygote
│   └── /opt/google/chrome/chrome --type=zygote
│       ├── /opt/google/chrome/chrome --type=zygote
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       │   └── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
│       └── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
```



How are processes created?

A process will **fork**(2) into a two new processes, which continue from the same place

The **parent** keeps the original PID and the **child** gets a new PID

Optionally, the new process (the child) **exec**(3) and begin running a new program



How many “henlo”s get printed?

```
int main( void ) {  
    fork();  
    printf( "henlo: %d\n", getpid() );  
}
```



How many “henlo”s get printed?

PID 1000

```
int main( void ) {  
    → fork();  
    printf( "henlo: %d\n", getpid() );  
}
```

PID 1000

```
int main( void ) {  
    fork();  
    → printf( "henlo: %d\n", getpid() );  
}
```

PID 2000

```
int main( void ) {  
    fork();  
    → printf( "henlo: %d\n", getpid() );  
}
```



How many “henlo”s get printed?

```
int main( void ) {  
    fork();    2^1  
    fork();    2^2  
    fork();    2^3  
    printf( "henlo: %d\n", getpid() );  
}
```

```
$ ./fork  
henlo 104632  
henlo 104635  
henlo 104634  
henlo 104633  
henlo 104638  
henlo 104639  
henlo 104637  
henlo 104636
```



Have you seen this code before?

```
:(){ :|:& };;:
```

```
bomb() {  
    bomb | bomb &  
} bomb
```



Making the child do something

```
int main( void ) {  
    if ( fork() > 0 ) {  
        /* parent process */  
        wait( NULL );  
    } else {  
        /* child process */  
        execv( "/bin/bash", NULL );  
    }  
}
```

```
/usr/bin/zsh  
\_ ./fork_exec  
    \_ [bash]
```



htop exploration

- ▣ htop(1): interactive terminal process manager
- ▣ <https://peteris.rocks/blog/htop/>
- ▣ PID/User/Command/CPU%/MEM%/TIME: self-explanatory
- ▣ PRI/NI: Priority and Niceness (take cs162, or google scheduling policy/SCHED_OTHER)
- ▣ VIRT/RES/SHR: virtual image, resident size, shared memory usage
- ▣ S: Status (S = sleep, R = running, T = terminated, Z = zombie)



What does this code do?


```
int main( void ) {  
    if ( fork() > 0 ) {  
        /* parent process */  
        sleep( 1 );  
    } else {  
        /* child process */  
        exit( 1 );  
    }  
}
```



What does this code do?

```
\_ /usr/bin/zsh
```

Parent  _ ./zombie-creator

Child  _ [zombie-creator] <defunct>




Zombie Process

When a child has died but
has not been “reaped”

Child metadata stays in
process table so parent can
collect exit status

Totally normal, all children
that exit are zombies!



How to kill zombie process



124

I launched my program in the foreground (a daemon program), and then I killed it with `kill -9`, but I get a zombie remaining and I m not able to kill it with `kill -9`. How to kill a zombie process?



share edit flag



50

edited Dec 10 '14 at 18:36



octosquidopus

1,125 ● 3 ● 20 ● 38

asked Jun 5 '13 at 16:14



MOHAMED

18.2k ● 28 ● 97 ● 169

5 Answers

active

oldest

votes



168

A zombie is already dead, so you cannot kill it.



An example of how you might want to signal to your process that it is the parent of a zombie child that this is achieved, make sure you kill processes that you do not intend to kill recursively using the set of signal handlers.

kill -9 \$(cat /dev/null & kill -9 \$?)

asked Jun 5 '13 at 16:14

answered Dec 10 '14 at 18:36



MOHAMED

18.2k ● 28 ● 97 ● 169

1

MOHAMED Jun 5 '13 at 16:14



What happens when a process dies

```
if ( fork() > 0 ) {  
    /* parent process */  
    sleep( 1 );  
} else {  
    /* child process */  
    exit( 1 );  
}
```

When a process exits,
returns int (exit code
0-255)

0 is success and
anything else indicates
an error



What happens when a process dies

```
if ( fork() > 0 ) {  
    /* parent process */  
    sleep( 1 );  
} else {  
    /* child process */  
    exit( 123 );  
}
```

**When a process exits,
returns int (exit code
0-255)**

**0 is success and
anything else indicates
an error**



Parents need to wait() on children

```
int main( void ) {  
    if ( fork() > 0 ) {  
        int status;  
        sleep( 1 );  
        wait( &status );  
        printf( "%d\n", WEXITSTATUS ( status ) );  
    } else {  
        exit( 123 );  
    }  
}
```


```
# keur @ magneto in ~/repos/temp  
$ ./good-parent  
123
```




What happens if the parent exits first?

```
int main( void ) {  
    if ( fork() > 0 ) {  
        /* parent process */  
        sleep ( 1 );  
        exit( 0 );  
    } else {  
        /* child process */  
        sleep( 100 );  
        exit( 123 );  
    }  
}
```

Process tree after start



Process tree after 1 second



What happens if the parent exits first?

If parent exits first, **orphan** processes are re-parented by the **init** process
init reaps all orphans that are zombies



When can zombies become a problem?

Parent doesn't `wait()` on children

Parent is long running process


Zombie child processes never become orphans

Resource leakage!



Zombie leakage in the wild

```
kitty
\_ /usr/bin/zsh
  \_ aerc
    \_ less -R
    \_ [awk] <defunct>
    \_ less -R
```

310bec27024579e7ada35585b3190ab875540804 – Kevin Kuehler  be4ea0d master

widgets/msgview: Reap the filter command

The filter command shells out and returns almost immediately. Call `Wait()` so the filter process gets reaped. Prior to this patch, `aerc` creates a zombie process for every email that is viewed.

Signed-off-by: Kevin Kuehler <keur@xcf.berkeley.edu>

```
    \_ less -R
    \_ [awk] <defunct>
    \_ less -R
    \_ [awk] <defunct>
    \_ less -R
    \_ [awk] <defunct>
    \_ less -R
    \_ [awk] <defunct>
```

Inter-process Communication

Various ways in which processes can communicate

- ▣ Exit codes
- ▣ Signals (e.g. SIGTERM, SIGKILL, SIGINT)
- ▣ Pipes (STDIN, STDOUT, STDERR)
- ▣ Sockets (UNIX socket, IP socket)
- ▣ Message Bus (e.g. dbus on Linux)
- ▣ ... and many many more...



Process Signals

- **SIGTERM**: tell a process to exit now
- **SIGKILL**: terminate process immediately
- **SIGINT**: interrupt, when you press Ctrl+C
- **SIGHUP**: the user closes the terminal window
- **SIGWINCH**: terminal window resized
- **SIGSTOP** / **SIGCONT**: stop/resume

SIGKILL and SIGSTOP cannot be handled



Services

What is a service?

- ▣ A service is a type of process known as a **Daemon**
 - ▣ A daemon is a noninteractive background process
 - ▣ Typically names end with a 'd', but not always the case
- ▣ Services are controlled by an **init** system
- ▣ Examples: sshd, httpd, rsyslogd, nginx, postfix, ...
- ▣ Not a strict definition

Services - What's the point?

- ▣ Can run for long periods of time
 - Useful in many cases, such as web servers
- ▣ Can be publicly accessible or shared between multiple users
 - Can be networked.
- ▣ Ex. sshd allows for incoming ssh connections - very important
- ▣ Can write your own services!



Systemd

- ▣ **Systemd** is an init system that manages processes and services
 - ▣ Most commonly used init system on modern Linux systems
- ▣ Provides tools for users to manage services
 - ▣ `systemctl` - start, stop, check status, and more
 - ▣ `journalctl` - check systemd journal
- ▣ Some debate in the free software community about systemd
 - ~~▣ systemd bad use runit~~



Systemd Unit Files

- ▣ Service behavior defined by systemd **unit files**
 - How should this service start up? How should it respond to various `systemctl` management commands?



A (simplified) Unit File: helloworld.service

[Unit]

Description=A simple unit file

Description - what the service does

[Service]

ExecStart=/usr/bin/helloworld

User=ocfstaff

Restart=always

Commands

[Install]

WantedBy=multi-user.target

When this unit should get started



Example Unit file - Nginx

```
[Unit]
Description=A high performance web server and a reverse proxy server
Documentation=man:nginx(8)
After=network.target

[Service]
Type=forking
PIDFile=/run/nginx.pid
ExecStartPre=/usr/sbin/nginx -t -q -g 'daemon on; master_process on;'
ExecStart=/usr/sbin/nginx -g 'daemon on; master_process on;'
ExecReload=/usr/sbin/nginx -g 'daemon on; master_process on;' -s reload
ExecStop=--/sbin/start-stop-daemon --quiet --stop --retry QUIT/5 --pidfile /run/nginx.pid
TimeoutStopSec=5
KillMode=mixed

[Install]
WantedBy=multi-user.target
```



systemctl

UNIT	LOAD	ACTIVE	SUB	DESCRIPTION
proc-sys-fs-binfmt_misc.automount	loaded	active	running	Arbitrary Executable File Formats File System Automount Point
sys-devices-pci0000:00-0000:00:03.0-virtio0-net-ens3.device	loaded	active	plugged	Virtio network device
sys-devices-pci0000:00-0000:00:05.0-virtio1-block-vda-vda1.device	loaded	active	plugged	/sys/devices/pci0000:00/0000:00:05.0/virtio1/block/vda/vda1
sys-devices-pci0000:00-0000:00:05.0-virtio1-block-vda-vda2.device	loaded	active	plugged	/sys/devices/pci0000:00/0000:00:05.0/virtio1/block/vda/vda2
sys-devices-pci0000:00-0000:00:05.0-virtio1-block-vda.device	loaded	active	plugged	/sys/devices/pci0000:00/0000:00:05.0/virtio1/block/vda
sys-devices-platform-serial8250-tty-ttyS1.device	loaded	active	plugged	/sys/devices/platform/serial8250/tty/ttyS1
sys-devices-platform-serial8250-tty-ttyS2.device	loaded	active	plugged	/sys/devices/platform/serial8250/tty/ttyS2
sys-devices-platform-serial8250-tty-ttyS3.device	loaded	active	plugged	/sys/devices/platform/serial8250/tty/ttyS3
sys-devices-pnp0-00:04-tty-ttyS0.device	loaded	active	plugged	/sys/devices/pnp0/00:04/tty/ttyS0
sys-subsystem-net-devices-ens3.device	loaded	active	plugged	Virtio network device
-.mount	loaded	active	mounted	Root Mount
dev-hugepages.mount	loaded	active	mounted	Huge Pages File System
dev-mqueue.mount	loaded	active	mounted	POSIX Message Queue File System
proc-sys-fs-binfmt_misc.mount	loaded	active	mounted	Arbitrary Executable File Formats File System
run-user-49390.mount	loaded	active	mounted	/run/user/49390
sys-kernel-debug.mount	loaded	active	mounted	Debug File System
systemd-ask-password-console.path	loaded	active	waiting	Dispatch Password Requests to Console Directory Watch
systemd-ask-password-wall.path	loaded	active	waiting	Forward Password Requests to Wall Directory Watch
init.scope	loaded	active	running	System and Service Manager
session-c15.scope	loaded	active	abandoned	Session c15 of user rrchan
session-c85.scope	loaded	active	running	Session c85 of user rrchan
• apache2.service	loaded	failed	failed	The Apache HTTP Server
console-setup.service	loaded	active	exited	Set console font and keymap
cpufrequtils.service	loaded	active	exited	LSB: set CPUFreq kernel parameters
cron.service	loaded	active	running	Regular background program processing daemon
dbus.service	loaded	active	running	D-Bus System Message Bus
getty@tty1.service	loaded	active	running	Getty on tty1
irqbalance.service	loaded	active	running	irqbalance daemon
keyboard-setup.service	loaded	active	exited	Set the console keyboard layout
kmmod-static-nodes.service	loaded	active	exited	Create list of required static device nodes for the current kernel
loadcpufreq.service	loaded	active	exited	LSB: Load kernel modules needed to enable cpufreq scaling
lvm2-lvmetad.service	loaded	active	running	LVM2 metadata daemon
lvm2-monitor.service	loaded	active	exited	Monitoring of LVM2 mirrors, snapshots etc. using dmeventd or progres
munin-node.service	loaded	active	running	Munin Node
netfilter-persistent.service	loaded	active	exited	netfilter persistent configuration
networking.service	loaded	active	exited	Raise network interfaces
nginx.service	loaded	active	running	A high performance web server and a reverse proxy server
node_exporter.service	loaded	active	running	Prometheus node_exporter
ntp.service	loaded	active	running	LSB: Start NTP daemon
php7.3-fpm.service	loaded	active	running	The PHP 7.3 FastCGI Process Manager
postfix.service	loaded	active	exited	Postfix Mail Transport Agent
postfix@.service	loaded	active	running	Postfix Mail Transport Agent (instance -)

`systemctl status [name]`

```
rrchan@solarflare:~$ systemctl status nginx
```

```
● nginx.service - A high performance web server and a reverse proxy server
   Loaded: loaded (/lib/systemd/system/nginx.service; enabled; vendor preset: enabled)
   Active: active (running) since Sat 2019-10-19 09:35:48 PDT; 2 days ago
     Docs: man:nginx(8)
  Main PID: 23491 (nginx)
    Tasks: 3 (limit: 4915)
   CGroup: /system.slice/nginx.service
           └─23491 nginx: master process /usr/sbin/nginx -g daemon on; master_process on;
              └─23494 nginx: worker process
                 └─23496 nginx: worker process
```



More systemctl!

- ▣ `systemctl start [name]` - starts a service
- ▣ `systemctl stop [name]` - stops a service
- ▣ `systemctl restart [name]` - restarts a service
- ▣ `systemctl reload [name]` - reload a service's configuration
- ▣ `systemctl enable [name]` - sets a service to start on boot
- ▣ `systemctl disable [name]` - opposite of *enable*



Example - sshd

- Handles ssh connections

PID	Command
1	/lib/systemd/systemd --system --deserialize 19
24854	└─ /usr/sbin/sshd -D
6392	└─ sshd: rrchan [priv]
6411	└─ sshd: rrchan@pts/0
6413	└─ -bash
14261	└─ htop

- Notice that *bash* and *htop* are children of *sshd*, since I am ssh'd into this machine.



Example - nginx and httpd

- nginx - http webserver daemon named 'Nginx'
- httpd - http webserver daemon named 'Apache'

Both examples of services which listen on port 80 and serve content.

```
• nginx.service - A high performance web server and a reverse proxy server
  Loaded: loaded (/lib/systemd/system/nginx.service; enabled; vendor preset: enabled)
  Active: active (running) since Sat 2019-10-19 09:35:48 PDT; 2 days ago
    Docs: man:nginx(8)
  Main PID: 23491 (nginx)
    Tasks: 3 (limit: 4915)
  CGroup: /system.slice/nginx.service
          └─23491 nginx: master process /usr/sbin/nginx -g daemon on; master_process on;
             └─23494 nginx: worker process
                └─23496 nginx: worker process
```



Bonus: Runit

- ▣ A better init system than systemd (bc I said so)
- ▣ Minimal
- ▣ Understand your services
- ▣ Average Runit User:



About Runit

- ▣ Created for the Void Linux distribution
 - ▣ Artix, Parabola
- ▣ Uses file, symbolic link, and folder management to start, stop, enable, disable, etc
- ▣ Everything is executable!



Hierarchy

- ▣ runsvdir: `/etc/runit/runsvdir/current` OR `/run/runit/service` #symlink
- ▣ Service Directory: `/etc/runit/sv/myservice`
 - ▣ `../myservice/run` #executable
 - ▣ `../myservice/log/run` #executable
 - ▣ And more!



Example: dbus

```
[adarsh@maxwell /etc/runit/sv]$ sudo tree dbus
dbus
├── check
├── log
│   └── run
│       └── supervise
│           ├── control
│           ├── lock
│           ├── ok
│           ├── pid
│           ├── stat
│           └── status
└── run
    └── supervise
        ├── control
        ├── lock
        ├── ok
        ├── pid
        ├── stat
        └── status
```

```
[adarsh@maxwell /etc/runit/sv]$ bat dbus/run
```

	File: dbus/run
1	#!/bin/sh
2	dbus-uuidgen --ensure=/etc/machine-id
3	[-d /run/dbus] install -m755 -g 81 -o 81 -d /run/dbus
4	exec dbus-daemon --system --nofork --nopidfile

```
[adarsh@maxwell /etc/runit/sv]$ bat dbus/log/run
```

	File: dbus/log/run
1	#!/bin/sh
2	exec 2>&1; set -e
3	
4	[-d /var/log/dbus] install -dm 755 /var/log/dbus
5	
6	exec svlogd -tt /var/log/dbus



Example: udevd

```
[adarsh@maxwell /etc/runit/sv]$ sudo tree udevd
udevd
├── finish
├── run
└── supervise
    ├── control
    ├── lock
    ├── ok
    ├── pid
    ├── stat
    └── status
```

```
[adarsh@maxwell /etc/runit/sv]$ sudo cat udevd/supervise/pid
1059
[adarsh@maxwell /etc/runit/sv]$ sudo cat udevd/supervise/stat
run
```

```
[adarsh@maxwell /etc/runit/sv]$ bat udevd/finish
```

File: udevd/finish	
1	#!/bin/sh
2	[-S /run/udev/control] && rm -f /run/udev/control



Service Management

- Enabling:

`ln -s /etc/runit/sv/dbus /run/runit/service/`

- Disabling:

`rm /run/runit/service`

- Stop on boot:

`touch /etc/runit/sv/dbus/down`

- Start on boot:

`rm /etc/runit/sv/dbus/down`



Service Management

- ▣ Start: `sv up dbus`
- ▣ Stop: `sv down dbus`
- ▣ Run Once: `sv once dbus`
- ▣ Restart: `sv status dbus`
- ▣ Status: `sv status dbus`
- ▣ Etc...!



Minimal Philosophy

- ▣ Always prefer minimal, simple software
 - ▣ Readable
 - ▣ Bug Free
 - ▣ Simple
 - ▣ Unix Philosophy:
Good at one thing and one thing only
- ▣ Also systemd is bad
- ▣ Like really bad



Like Really Bad (not OCF endorsed msg)

- ▣ <https://artixlinux.org/faq.php>
- ▣ <https://chiefio.wordpress.com/2016/05/18/systemd-it-keeps-getting-worse/>
- ▣ https://www.theregister.com/2019/01/31/systemd_exploit/

