

## Lecture 5

ncograin

(slide credits abizer, longlian, night)

# Networking 101

---

# Course Resources

---

- ▣ Your facilitators!
- ▣ Ed, Gradescope
- ▣ OCF Slack ([ocf.io/slack](https://ocf.io/slack)) or Discord ([ocf.io/discord](https://ocf.io/discord))  
#decal-general
- ▣ All materials available at [decal.ocf.io](https://decal.ocf.io)
- ▣ Ask questions / work on lab with us during lab sessions!  
(Tuesday 8-9pm in OCF Lab)



# Outline

---

- What makes up a network system?
- Addressing & Routing
- Protocols
- SysAdmin Tools



**What makes up a  
network system?**

# Big Question: Why do we need networks?

---

- ▣ We want easy communication!

What do we need for easy communication?

- ▣ Low latency
- ▣ High bandwidth
- ▣ Ability to easily identify a computer
- ▣ Fault-tolerance



# Conceptual Model: TCP/IP

---

- ▣ 4-layer model
- ▣ Loosely defined (functions can be in-between levels)
- ▣ Higher layers abstract lower layers

Application

Transport

Internet

Link

# Link Layer

---

- ▣ Communication between physically connected nodes
- ▣ MAC Addresses
- ▣ ARP/NDP
- ▣ Underlying hardware implementations are abstracted outside of TCP/IP model
- ▣ “Who can I talk to without going through a router?”



# Internet Layer

---

- ▣ Transmitting data to different networks
- ▣ Global Routing (IP Addresses)
- ▣ No guarantees on reliability (just like real life!)





# Transport Layer

---

- ▣ Defines protocols for communication between computers
- ▣ Dealing with unreliability of sending packets
- ▣ TCP/UDP



# Application Layer

---

- ▣ Application-specific, actually deals with interpreting the data being sent
- ▣ Most “user-facing” functions live here
- ▣ File sharing, message passing, database access
- ▣ Examples: HTTP/FTP/DNS



# Addressing & Routing

# Media Access Control (MAC) Addresses

---

- ▣ Looks like this: 00:14:22:01:2a:c5
- ▣ 48 bits, divided into 6 octets
- ▣ First 3 octets identify the interface manufacturer
- ▣ Unique to each device
- ▣ Used to identify devices on local network
- ▣ How do we talk to devices outside the local network?

# IP Addresses (IPv4)

---

- ▣ IP addresses uniquely identify a host
- ▣ IP addresses can be changed
- ▣ 32 bits long, divided into 4 octets, look like:  
123.123.123.123
- ▣ IPv4 octets are commonly expressed in decimal notation, although alternative formats, though rarely employed, are also permissible.



# A note on bits, decimal and hexadecimal

---

- ▣ Different ways to write a number!
  - Decimal: uses 0 through 9
    - Add another digit after reaching 9
  - Hex: 0 through 9, a through f
    - Add another digit after reaching f
  - Binary: uses bits (0 or 1)
    - Add another digit... every two numbers... :(
- ▣ . or : in addresses are just separators. Make it easy to read



# Example: Sixty

---

Decimal:

60

Hex:

3c

(often written 0x3c)

Binary:

0b111100

It's more convenient  
to use hex instead!

Used in IPv4

Used in MAC  
addresses, IPv6



# Subnets

---

- ▣ Remember that IP addresses are represented with binary
- ▣ Sometimes, you'll see 192.168.1.0/24
- ▣ The /24 is a subnet mask
- ▣ Subnet: a range of IP addresses (e.g. 192.168.1.0-255)
- ▣ The first 24 bits of the address: network prefix
- ▣ The last 8 bits: host identifier





# IPv6

---

- ▣ There are only 4.3 billion IPv4 addresses, but almost 8 billion people in the world - what do we do?
- ▣ IPv6 addresses contain 128 bits
- ▣ Formatted as follows: 1234:5678:89:0:ab:cd:ef:beef
- ▣ We have a million billion IPv6 addresses for each cell in every human on the planet (enough?)



# Address Resolution Protocol (ARP)

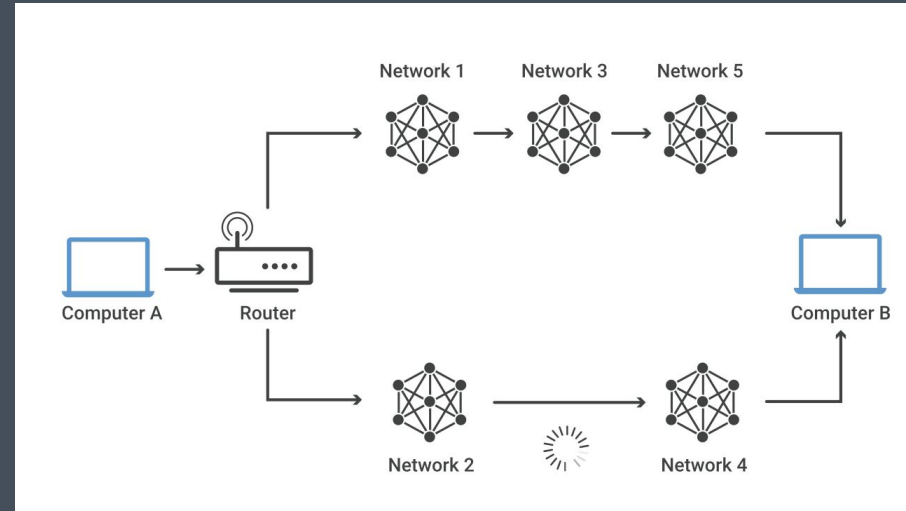
---

- ▣ Translates IP address to a specific MAC address
- ▣ ARP Request: asks devices on local network who an IP address belongs to
- ▣ Kernel caches values in an “ARP table”



# Routing

- How do we decide where to jump to?
- Routers maintain routing tables, which tell us where to hop to next



169.229.226.0/24	ocf network
169.229.0.0/16	berkeley network
128.32.0.0/16	berkeley network
0.0.0.0/0	via 128.32.0.39

OCF router routing table

# Domain Name System (DNS)

---

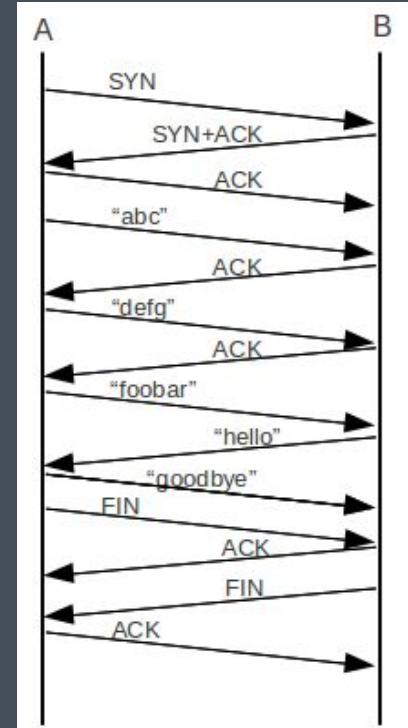
- ▣ Translates domains (as found in URLs) to IP addresses
- ▣ Can manually set domain -> IP mappings, or can use external DNS resolver



# Protocols

# Transmission Control Protocol (TCP)

- ▣ Ensures reliable transmission
- ▣ Connection-oriented: need to initiate a connection before sending data
- ▣ Used when reliability is more important than speed
- ▣ Examples: website loading, file transfer



# User Datagram Protocol (UDP)

---

- ▣ No reliability guarantees
- ▣ Does not establish a connection
- ▣ Good for real-time applications where some loss is acceptable
- ▣ Examples: (lossy) streaming, gaming



# Internet Control Message Protocol (ICMP)

---

- ▣ Not used to transmit data
- ▣ Technically not a “transport protocol”
- ▣ Used to transmit error messages and status info
- ▣ Used by diagnostic tools





# Internet Control Message Protocol (ICMP)

Type	Code	Status	Description
0 – Echo Reply <sup>[6]:14</sup>	0		Echo reply (used to ping)
3 – Destination Unreachable <sup>[6]:4[9]</sup>	0		Destination network unreachable
	1		Destination host unreachable
	3		Destination port unreachable
8 – Echo Request	0		Echo request (used to ping)



# SysAdmin Tools

# ping

---

```
|
```

# ping

---

```
@ ~ > ping -c 5 google.com
PING google.com(sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e)) 56 data bytes
64 bytes from sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e): icmp_seq=1 ttl=114 time=3.15 ms
64 bytes from sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e): icmp_seq=2 ttl=114 time=3.13 ms
64 bytes from sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e): icmp_seq=3 ttl=114 time=3.18 ms
64 bytes from sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e): icmp_seq=4 ttl=114 time=3.11 ms
64 bytes from sfo07s26-in-x0e.1e100.net (2607:f8b0:4005:802::200e): icmp_seq=5 ttl=114 time=3.16 ms

--- google.com ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4007ms
rtt min/avg/max/mdev = 3.111/3.144/3.176/0.022 ms
```



# traceroute

---

```
|
```

# traceroute

```
@ ~ > traceroute google.com                                     lmathias@tsunami
traceroute to google.com (142.250.189.206), 30 hops max, 60 byte packets
 1  vlan635.inr-350-reccev.berkeley.edu (169.229.226.1)  0.552 ms  0.511 ms  0.457 ms
 2  reccev-cev-cr1--xe-2-0-9.net.berkeley.edu (128.32.0.168)  1.445 ms  1.277 ms  1.278 ms
 3  sut-mdc-sr4--irb-117.net.berkeley.edu (128.32.82.177)  0.949 ms  0.823 ms  0.760 ms
 4  firewall--ethernet1-21-595.ocf.berkeley.edu (128.32.82.178)  0.824 ms  0.856 ms  0.813 ms
 5  sut-mdc-sr5--irb-116.net.berkeley.edu (128.32.82.173)  1.184 ms  1.149 ms  1.117 ms
 6  sut-mdc-cr1--et-0-1-1.net.berkeley.edu (128.32.255.40)  1.751 ms  reccev-cev-cr1--et-0-1-1.net.berkeley.edu (128.32.255.42)
   ) 1.796 ms  sut-mdc-cr1--et-0-1-1.net.berkeley.edu (128.32.255.40)  1.743 ms
 7  reccev-cev-br1--et-1-1-0.net.berkeley.edu (128.32.0.36)  1.710 ms  1.533 ms  1.466 ms
 8  oak-agg8--ucb--100g.cenic.net (137.164.3.26)  3.319 ms  3.288 ms  2.413 ms
 9  74.125.48.172 (74.125.48.172)  2.566 ms  142.250.164.8 (142.250.164.8)  2.732 ms  2.691 ms
10  108.170.243.1 (108.170.243.1)  4.148 ms  108.170.242.225 (108.170.242.225)  2.792 ms  108.170.243.1 (108.170.243.1)  4.071
   ms
11  142.251.224.175 (142.251.224.175)  2.837 ms  142.251.224.173 (142.251.224.173)  2.683 ms  142.251.224.175 (142.251.224.175)
   3.104 ms
12  sfo03s25-in-f14.1e100.net (142.250.189.206)  2.889 ms  2.857 ms  2.848 ms
```



# arp

```
~ > arp
```

Address	HWtype	HWaddress	Flags	Mask	Interface
dhcp-169-229-226-182.OC	ether	7e:e7:6d:6e:4d:cd	C		ens3
supernova.OCF.Berkeley.	ether	52:54:00:c9:60:55	C		ens3
anthrax.OCF.Berkeley.ED	ether	52:54:00:c3:f9:d8	C		ens3
dataloss.OCF.Berkeley.E	ether	ac:1f:6b:1a:8e:92	C		ens3
dementors.OCF.Berkeley.	ether	52:54:00:a6:16:f6	C		ens3
dhcp-169-229-226-178.OC	ether	6e:74:84:51:0f:d8	C		ens3
bigbang.OCF.Berkeley.ED	ether	d0:17:c2:d2:08:62	C		ens3
blackhole.OCF.Berkeley.	ether	00:1c:73:3d:2f:c2	C		ens3
dhcp-169-229-226-143.OC	ether	ba:fc:93:84:3c:61	C		ens3
vlan635.inr-350-reccev.	ether	88:e6:4b:0d:cc:6e	C		ens3
dhcp-169-229-226-142.OC	ether	bc:d0:74:12:a3:e2	C		ens3
dhcp-169-229-226-171.OC	ether	04:ea:56:de:e3:56	C		ens3
dhcp-169-229-226-145.OC	ether	3c:22:fb:0f:d1:22	C		ens3
flood.OCF.Berkeley.EDU	ether	52:54:00:34:09:e6	C		ens3



# dig





# dig

---

```
@ ~ > dig google.com

; <<>> DiG 9.16.37-Debian <<>> google.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 20879
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:: udp: 1232
; COOKIE: 7b5c9f7901f27f540100000063feda0345fb6daf74a8b74b (good)
;; QUESTION SECTION:
;google.com.                IN      A

;; ANSWER SECTION:
google.com.                 202     IN      A      142.251.46.174

;; Query time: 0 msec
;; SERVER: 2607:f140:8801::1:22#53(2607:f140:8801::1:22)
;; WHEN: Tue Feb 28 20:52:19 PST 2023
;; MSG SIZE rcvd: 83
```



# ip

---

Can do many things, including displaying/modifying ip addresses and routing info. Example: ip addr

```
~ > ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq state UP group default qlen 1000
    link/ether 52:54:00:c9:60:55 brd ff:ff:ff:ff:ff:ff
    inet 169.229.226.36/24 brd 169.229.226.255 scope global ens3
        valid_lft forever preferred_lft forever
    inet6 2607:f140:8801::1:36/64 scope global
        valid_lft forever preferred_lft forever
    inet6 fe80::5054:ff:fec9:6055/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:48:6d:20:27 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
```

