

Lecture 11

jaysa

(slide credits fy dai)

Config Management (with Puppet)

Congratulations, you're hired!


You've earned your Open
Computing Facility
Full-Fledged Linux Expert
Certification™*, and now
you have been hired to
manage a computer lab...

Problem 1: Managing a fleet

Suppose you have a bunch of computers... and you decide that everybody computer in the lab needs Minecraft installed.

- ▣ Manually install on each desktop?
- ▣ Write a script to SSH into each desktop and install?

How do I deploy updates to a fleet of existing computers?



Problem 2: Provisioning new machines

Suppose your computer lab buys a new computer... and you want to set it up the same as the others

- ▣ Remembering to install each software and configure it correctly is pretty difficult!
 - ▣ Eg: Have Minecraft use more than 1Gb RAM by default

How do I provision new machines?




Problem 3: Communicating changes

Suppose you are running a Minecraft server... and you realize some change you made long ago broke a thing.

- ▣ How do you figure out what settings you had changed around that time?

How do you communicate changes in software configuration to future you (and others)?



Config Management!

✓ Add minecraft-launcher to packages.pp

Browse files

master

fydai committed on Feb 12 Verified 1 parent [be4bc6a](#) commit [466fc20ef26470425792c5d9df4bc03ba21893c8](#)

Showing 1 changed file with 1 addition and 1 deletion.

Unified Split

2 modules/ocf_desktop/manifests/packages.pp

29	29	# FUSE
30	30	['fuse', 'exfat-fuse'];;
31	31	# games
32	-	['armagetronad', 'gl-117', 'gnome-games', 'wesnoth', 'wesnoth-music'];;
32	+	['armagetronad', 'gl-117', 'gnome-games', 'minecraft-launcher', 'wesnoth', 'wesnoth-music'];;
33	33	# graphics/plotting
34	34	['r-cran-rgl', 'jupyter-qtconsole', 'rstudio'];;
35	35	# input method editors



Config Management

Config management software makes it as easy as possible to bootstrap new machines, configure running software, and allows configuration to be stored as code “configuration as code” philosophy

- Solves problem 1 (updating computers) by having an unified update mechanism
- Solves problem 2 (bootstrapping new computers) by having all the changes necessary in some centralized repository
- Solves problem 3 (communication) by allowing you to use standard development practices (mainly git) to record your changes, and communicate with others



Configuration Management Philosophies

Imperative

- Write the **set of tasks**"
- "Install minecraft", then "add a line to the config file", then "run minecraft"

Declarative

- Write the **desired state**
- "Ensure minecraft is installed, the config file has line <X> in it, and ensure that minecraft is running"



Configuration Management Software

from Stack Overflow

Developer Survey 2023:

% of respondents who use:

- Kubernetes - 19.02%
- Terraform - 11.3%
- Ansible - 8.62%
- Puppet - 1.12%
- Chef - 0.94%





Let's learn

puppet !

Puppet

- ▣ Released 2005; has open source and enterprise closed source version
- ▣ Declarative philosophy, with some Imperative components when necessary
- ▣ Originally built on Ruby, now its own configuration language
- ▣ **“Pull model”** - Configured machines ask for an update
 - ▣ At OCF, puppet is scheduled to run every 30 minutes

Used at places like

- ▣ OCF
(github.com/ocf/puppet)
- ▣ CS 162
(github.com/Berkeley-CS-162/vagrant/tree/master/modules/cs162)
- ▣ Wikimedia
(github.com/wikimedia/puppet)
- ▣ Github
- ▣ Lyft



What happens when Puppet is run?

1. **Client** asks server for an update
2. **Server** asks client for a list of Facts
3. **Client** responds with the facts
4. **Server** responds with configuration
5. **Client** makes the necessary changes to ensure its current configuration matches the configuration given by the server

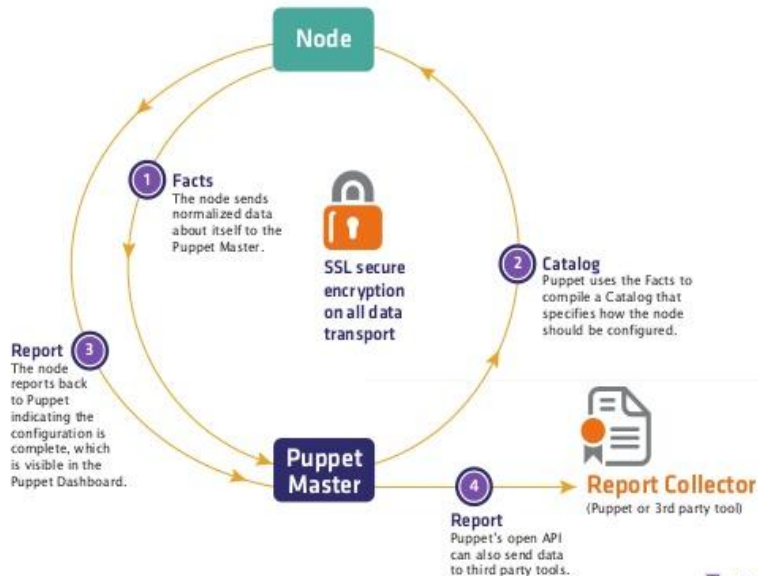
1. "I want to be configured as a Minecraft Server"
2. "Ok, send me your hostname, and RAM"
3. "My hostname is zombies.ocf.berkeley.edu and I have 4GB RAM"
4. "Ensure the Minecraft server is running, with hostname zombies.ocf.berkeley.edu, 4GB RAM, with this configuration file"
5. "The minecraft server is currently running, but the configuration file has been updated, I will fetch the updated version"



What happens when Puppet is run?

1. **Client** = Node
2. **Server** = Puppet Master

Lifecycle of a Puppet Run



Puppet Code

- Mostly, we have:
 - **Files** - contains static files
 - **Templates** - contain templates (Ruby style)
 - **Manifests** - the heart of the configuration, specifies the desired states
- Other sections that are occasionally used
 - **Facts** - Ways to extract data needed for configuration
 - **Functions** - if you need extra something fancy data structure manipulating
- Dependencies need to be explicitly described
 - Puppet is allowed to run code in any order that it sees fit
 - If you have code installing Minecraft, and running Minecraft, you need to tell puppet to install Minecraft before running it



Example Puppet Code 1

```
user { 'ocftv':  
  comment => 'TV NUC',  
  home    => '/opt/tv',  
  groups  => ['sys', 'audio'],  
  shell   => '/bin/bash';  
}  
  
file {  
  # Create home directory for ocftv user  
  '/opt/tv':  
    ensure => directory,  
    owner  => ocftv,  
    group  => ocftv,  
    require => User['ocftv'];  
}
```

What does this do?

A: Adds a user (ocftv)
and a home directory
(/opt/tv)



Example Puppet Code 2

```
package { 'nginx';; }
service { 'nginx':
  require => Package['nginx'],
  subscribe => Class['ocf::ssl::default'],
}

file {
  '/etc/nginx/conf.d/local.conf':
    content => template('ocf_apphost/local.conf.erb'),
    require => Package['nginx'],
    notify => Service['nginx'];
}

1 # raise the hash bucket size for server names since we use really long server
2 # names (like something.apphost.ocf.berkeley.edu)
3 #
4 # http://nginx.org/en/docs/http/server_names.html
5 server_names_hash_bucket_size 128;
6
7 ssl_dhparam /etc/ssl/dhparam.pem;
8 ssl_protocols <%= @ssl_protocols %>;
9 ssl_ciphers '<%= @ssl_ciphersuite %>';
10
11 # combined log format, with virtual host added (rt#4459)
12 log_format vhost '$host $remote_addr - $remote_user [$time_local] '
13                 '$request' $status $body_bytes_sent '
14                 '$http_referer' '$http_user_agent';
15
16 # increase client max request body size (default is 1MiB)
17 client_max_body_size 20M;
```


What does this do?

A: Runs a web server!



Puppet at the OCF (1/2)

github.com/ocf/puppet



- ▣ Originally started in 2012, from the “Configuration is edited directly on the server, and desktops manually” model
 - ▣ We only had 10 desktops so this was kinda okay
- ▣ 11 years later, all of the OCF’s machines* runs off the puppet repository
 - ▣ Desktops
 - ▣ Thing behind the TV
 - ▣ Hypervisors (things running the VMs)
 - ▣ VMs (Running all the Networked Services you learned about)
 - Including the puppet server itself

* actually in 2024 we moved a bunch of things to nix and plan to move the rest by 2025



Puppet at the OCF (2/2)

- ▣ All the code is split into modules
 - ocf_tv
 - ocf_desktop
 - ocf_www
 - ocf_printhost
- ▣ Common OCF modules for shared configuration
 - ocf::ssl for (I need a web certificate)
 - ocf::auth for LDAP and Kerberos and sudoers configuration

Bonus Slide: Terraform



- Has integrated APIs to provision machines declaratively on cloud platforms
- This is the code used to generate your decal VMs!
- The alternative (which we seriously considered) was clicking “New droplet” 80 times.
- If you’re interested in seeing more of the decal VM Terraform code, here’s the repository:

github.com/0xcf/decal-automation

```
54     provisioner "remote-exec" {
55         connection {
56             type      = "ssh"
57             user       = "root"
58             private_key = "${file("${var.pvt_key_file}")}"
59             host        = self.ipv4_address
60         }
61         inline = [
62             # Set the hostname to be FQDN
63             "sudo hostname ${each.value.username}.decal.xcf.sh",
64
65             # Add the user and give them root access
66             "sudo useradd ${each.value.username} -s /bin/bash -m",
67             "sudo echo ${each.value.username}:${each.value.password} | sudo chpasswd",
68             "sudo usermod -aG sudo ${each.value.username}",
69             "sudo service sshd restart",
70             "sudo passwd -e ${each.value.username}",
71
72             # Populate the motd with data
73             "sudo sed -i 's/${HOSTNAME}/${each.value.username}/g' /etc/motd",
74             "sudo sed -i 's/${IP}/${self.ipv4_address}/g' /etc/motd",
75
76             # Reboot to allow MOTD to change. This is a hack.
77             "sudo shutdown -r +60"
78         ]
79     }
```

